

egk-mixnet-maths DRAFT

John Caron

1/22/2024

This is the mathematical description of the code in the egk-mixnet library [12], which is an extraction of the Terelius / Wikström (TW) mixnet algorithm [1] [2] implementation from the Verificatum library [4].

The math here mostly recapitulates the work of Wikström [3]; Haenni et. al. [5] [6] that explain the Terelius / Wikström (TW) mixnet algorithm; and the work of Haines [8] that gives a formal proof of security of TW when the shuffle involves vectors of ciphertexts.

The verification equations are well documented in [3]. The proof equations are reverse engineered from reading the Verificatum code, and are apparently not otherwise documented, in particular not in [3], although likely they are implied in [1] but using different notation.

The proof equations are implemented in the egk-mixnet library class `ShuffleProver` and the verifier equations in `ShuffleVerifier`. The `ShuffleVerifier` has also been tested against the proofs output by Verificatum itself, leading to some confidence that these equations capture the TW algorithm as implemented in Verificatum.

Instead of providing pseudo-code, the Kotlin library code is the implementation of the math described here. The code implementing the algorithm is mostly separate from the workflow and the serialization code. It can act as a reference and comparison for ports to other languages and to other applications needing a mixnet.

Contents

1	Definitions	3
1.1	The ElectionGuard Group	3
1.2	ElGamal Encryption and Reencryption	3
1.3	Permutations	4
1.3.1	Permutation Commitments	4
2	Terelius / Wikström Algorithm	5
2.1	Definitions	5
2.2	Non-interactive Calculations	6
2.2.1	Generators	6
2.2.2	BatchingVector and Challenge	6
2.3	Shuffle	7
2.4	Proof	7
2.4.1	Commitment to permutation	7
2.4.2	Commitment to shuffle	7
2.4.3	Reply to challenge v:	9
2.4.4	The ProofOfShuffle Data Structure	9
2.5	Proof Verification	10
3	Performance	12
3.1	Operation counts	12
3.2	Regular vs accelerated exponentiation time	13
3.3	Timing results	14
3.4	Discussion of Timing results	16
3.4.1	Parallelization	16
3.4.2	Algorithm differences	17
A	ElGamal Encryption and Reencryption	18
B	Alternative Calculation of B and B'	19
C	Implementation notes	21
	References	21

1 Definitions

The ElectionGuard Kotlin library [10] and ElectionGuard 2.0 specification [9] are used for the cryptography primitives, in particular the parameters for Z_p^r , the variant of ElGamal exponential encryption, and the use of HMAC-SHA-256 for hashing.

1.1 The ElectionGuard Group

- $Z = \{\dots, -3, -2, -1, 0, 1, 2, 3, \dots\}$ is the set of integers.
- $Z_n = \{0, 1, 2, \dots, n-1\}$ is the ring of integers modulo n .
- Z_n^* is the multiplicative subgroup of Z_n that consists of all invertible elements modulo n . When p is a prime, $Z_p^* = \{1, 2, 3, \dots, p-1\}$
- Z_p^r is the set of r -th-residues in Z_p^* . Formally, $Z_p^r = \{y \in Z_p^* \text{ for which there exists } x \in Z_p^* \text{ where } y = x^r \bmod p\}$. When p is a prime for which $p-1 = q \cdot r$ with q a prime that is not a divisor of the integer r , then Z_p^r is an order- q cyclic subgroup of Z_p^* , and for any $y \in Z_p^*$, $y \in Z_p^r$ if and only if $y^q \bmod p = 1$.
- We use $r = (p-1)/q$ to denote the cofactor of q , and choose a generator g of Z_p^r . See [9] for details.

1.2 ElGamal Encryption and Reencryption

A variant of exponential ElGamal is used for encryption and reencryption (also see [11] and Appendix A), using the group generator g , and a public key K :

$$Encr(m, \xi) = (g^\xi, K^{m+\xi}) \quad (1)$$

$$Encr(0, \xi') = (g^{\xi'}, K^{\xi'}) \quad (2)$$

$$ReEncr(m, r) = (g^{\xi+r}, K^{m+\xi+r}) = Encr(0, r) * Encr(m, \xi) \quad (3)$$

1.3 Permutations

A *permutation* is a bijective map $\psi : 1..n \rightarrow 1..n$. If \vec{x} is a vector, $\psi(\vec{x})$ is the permutation of its elements, which we denote \vec{x}' , so that $x_i = x'_j$, where $i = \psi(j)$ and $j = \psi^{-1}(i)$. If M is a matrix, $\psi(M)$ is the permutation of its row vectors, to form a new matrix M' of the same shape.

A permutation ψ has a *permutation matrix* $B_\psi = (b_{ij})$, where $b_{ij} = 1$ if $\psi(i) = j$, otherwise 0, and where $\psi(\vec{x}) = B_\psi \vec{x}$ (matrix multiply).

If $B_\psi = (b_{ij})$ is an n -by- n matrix and $\vec{x} = (x_1, \dots, x_n)$ any vector of N independent variables, then B_ψ is a permutation matrix if and only if

$$\sum_{i=1}^n b_{ij} = 1 \quad (4)$$

$$\sum_{i=1}^n \sum_{j=1}^n b_{ij} x_i = \sum_{i=1}^n x_i \quad (5)$$

1.3.1 Permutation Commitments

For a vector $\vec{m} = (m_1..m_w) \in Z_q$, the *Pedersen commitment* to \vec{m} is

$$Commit(\vec{m}, cr) = g^{cr} * h_1^{m_1} * h_2^{m_2} * .. h_w^{m_w} = g^{cr} * \prod_{i=1}^w h_i^{m_i}$$

where \vec{h} are independent generators of Z_p^r and $cr \in Z_q$ is a randomization nonce.

If \vec{b}_j is the j^{th} column vector of a permutation matrix B_ψ , then the *permutation commitment to ψ* is defined as the vector of commitments to its columns:

$$Commit(\psi, \vec{cr}) = (Commit(\vec{b}_1, cr_1), .. Commit(\vec{b}_n, cr_w)) \quad (6)$$

$$Commit(\vec{b}_j, cr_j) = g^{cr_j} * \prod_{i=1}^w h_i^{b_{ij}} = g^{cr_j} h_i, \text{ for } i = \psi^{-1}(j) \quad (7)$$

$$Commit(\psi, \vec{cr}) = \{g^{cr_j} h_{\psi^{-1}(j)}\}, j=1..n \quad (8)$$

2 Terelius / Wikström Algorithm

2.1 Definitions

Let

- n = number of rows (eg ballots)
- width = number of ciphertexts in each row
- W = matrix of ciphertexts ($n \times \text{width}$), with entries $w_{i,j}$; its row vectors of width ciphertexts are $\vec{w}_i, i = 1..n$; and its column vectors of n ciphertexts are $\vec{w}_j, j = 1..\text{width}$
- R = matrix of reencryption nonces, where R_{ij} is the reencryption nonce for W_{ij} .
- W' is the shuffle of W = matrix of shuffled and reencrypted ciphertexts. Its entries, row vectors and column vectors are $w'_{i,j}, \vec{w}'_i, \vec{w}'_j$ respectively
- ψ = permutation function
- ψ^{-1} = inverse permutation function
- \vec{h} = independent generators of Z_p^r
- $h_0 = \vec{h}_1$
- \vec{e} = batching vector.
- $\vec{e}' = \psi(\vec{e})$ = permuted batching vector.
- $v \in Z_q$ = challenge.

We use one-based array indexing here for notational simplicity. The code itself uses zero based indexing.

2.2 Non-interactive Calculations

In the non-interactive version of the proof, the generators \vec{h} , batching vector \vec{e} , and challenge v must be generated with a deterministic calculation that both the Prover and the Verifier can independently make, to prevent them from being carefully chosen by a corrupt Prover to subvert the proof.

Generally these are done by Verificatum [4] in ways that are highly dependent on the Verificatum implementation. While it is possible to replicate these calculations, it is likely acceptable to substitute other calculations. Here we outline the current implementation in egk-mixnet, but these need to be reviewed by competent cryptographers and may need to be modified.

2.2.1 Generators

- Use the electionguard *parameterBaseHash()* on the group parameters, and an arbitrary string *mixName* to generate a seed.
- Use the electionguard *Nonces* class to create a vector of nonces from the seed.
- Calculate $h_1 = g^{\text{nonce}_1}$
- Calculate $h_i = h_0^{\text{nonce}_i}$

See *org.cryptobiotic.mixnet.getGeneratorsVmn()* in [12].

2.2.2 BatchingVector and Challenge

- Create a seed by hashing all of the following:
 - the electionguard *parameterBaseHash()* on the group parameters
 - the generators \vec{h}
 - the permutation commitments \vec{u}
 - the election public key K
 - all the ciphertexts from W
 - all the ciphertexts from W'

- Use the electionguard *Nonces* class to create a vector of nonces from the seed.
- Calculate \vec{e} , where $e_i = \text{nonce}_i$
- Calculate v by hashing the seed and an arbitrary string *mixName*.

See *org.cryptobiotic.mixnet.getBatchingVectorAndChallenge()* in [12].

2.3 Shuffle

Let R be a matrix of ($n \times \text{width}$) randomly chosen reencryption nonces. For each ciphertext in W :

$$\begin{aligned} \text{Reencrypt}(w_{ij}) &= \text{Reencrypt}(w_{ij}, r_{ij}) = \text{Encr}(0, r_{ij}) * w_{ij} \\ &= \text{Reencrypt}(W, R) \end{aligned}$$

Then W' is the permutation of the rows of $\text{Reencrypt}(W, R)$:

$$W' = \psi(\text{Reencrypt}(W, R))$$

The shuffle step chooses R and ψ , so $\text{Shuffle}(W) \rightarrow (R, \psi, W')$

2.4 Proof

2.4.1 Commitment to permutation

Choose a vector of random permutation nonces $\vec{c\vec{r}}$. Form public permutation commitments \vec{u} :

$$u_j = g^{c\vec{r}_j} h_i \quad \text{for } i = \psi^{-1}(j), j=1..n \quad (9)$$

2.4.2 Commitment to shuffle

Choose vectors of n random nonces $\vec{b}, \vec{\beta}, e\vec{p}s$ and random nonces α, γ, δ , all $\in Z_q$.

Form the following values $\in Z_p^r$:

$$A' = g^\alpha \prod_{i=1}^n h_i^{eps_i} \quad (10)$$

$$B_1 = g^{b_1} h_0^{e'_1}, \quad B_i = g^{b_i} (B_{i-1})^{e'_i}, \quad i = 2..n \quad (11)$$

$$B'_1 = g^{\beta_1} h_0^{eps_1}, \quad B'_i = g^{\beta_i} (B'_{i-1})^{eps_i}, \quad i = 2..n \quad (12)$$

$$C' = g^\gamma \quad (13)$$

$$D' = g^\delta \quad (14)$$

Also see Appendix B for a variation on computing B and B' .

Commitment to exponents

Choose *width* random nonces $\vec{\phi}$.

Form the following ciphertext values:

$$F'_j = Encr(0, -\phi_j) \cdot \prod_{i=1}^n (w'_{i,j})^{eps_i} \text{ for } j=1..\text{width} \quad (15)$$

Note that \vec{F}' has *width* components, one for each of the column vectors of $W' = \vec{w}'_j$. For each column vector, form the component-wise product of it exponentiated with $e\vec{p}s$. We can use any of the following notations to indicate this:

$$\begin{aligned} &= \prod_{i=1}^n (w'_{i,j})^{eps_i} \text{ } j=1..\text{width} \\ &= \prod_{i=1}^n (\vec{w}'_j)_i^{eps_i} \text{ } j=1..\text{width} \\ &= \prod_{i=1}^n (W')^{eps} \end{aligned}$$

This disambiguates the equations in Algorithm 19 of [3], for example: $\prod w_i^{e_i}$ and $\prod (w'_i)^{k_{E,i}}$.

2.4.3 Reply to challenge v:

A challenge $v \in Z_q$ is created as in 2.2.2, and the following values $\in Z_q$ are made as reply:

$$k_A = v \cdot \langle \vec{c}, \vec{e} \rangle + \alpha \quad (16)$$

$$\vec{k}_B = v \cdot \vec{b} + \vec{\beta} \quad (17)$$

$$k_C = v \cdot \sum_{i=1}^n cr_i + \gamma \quad (18)$$

$$k_D = v \cdot d + \delta \quad (19)$$

$$\vec{k}_E = v \cdot \vec{e'} + e\vec{p}s \quad (20)$$

and, with $\vec{R}_j = j$ th column of reencryption nonces R:

$$k_{F,j} = v \cdot \langle \vec{R}_j, \vec{e'} \rangle + \phi_j \text{ for } j=1..\text{width} \quad (21)$$

where \langle, \rangle is the inner product of two vectors, and \cdot is scalar multiply.

2.4.4 The ProofOfShuffle Data Structure

$\text{ShuffleProver}(R, \psi, W') \rightarrow \text{ProofOfShuffle}$

```
data class ProofOfShuffle(
  val mixName: String,
  val u: VectorP, // permutation commitment

  // Commitment of the Fiat-Shamir proof.
  val Ap: ElementModP,
  val B: VectorP,
  val Bp: VectorP,
  val Cp: ElementModP,
  val Dp: ElementModP,
  val Fp: VectorCiphertext, // size width

  // Reply of the Fiat-Shamir proof.
  val kA: ElementModQ,
  val kB: VectorQ,
  val kC: ElementModQ,
```

```

    val kD: ElementModQ,
    val kE: VectorQ,
    val kF: VectorQ, // size width
)

```

2.5 Proof Verification

The following equations are taken from Algorithm 19 of [3] and checked against the Verificatum implementation. The main ambiguity is in the meaning of $\prod_{i=1}^n w_i^{e_i}$ and $\prod_{i=1}^n (w'_i)^{k_{E,i}}$ in steps 3 and 5. These are interpreted as a short hand for *width* equations on the column vectors of W and W' , respectively, as detailed in *commitment to exponents* section above.

The Verifier is provided with:

- W = rows of ciphertexts (n x width)
- W' = shuffled and reencrypted rows of ciphertexts (n x width)
- the ProofOfShuffle

The \vec{h} (generators), \vec{e} nonces, and challenge v are deterministically recalculated from the algorithms described in section 2.2.

The following values $\in Z_p^r$ are computed:

$$A = \prod_{i=1}^n u_i^{e_i} \tag{22}$$

$$C = (\prod_{i=1}^n u_i) / (\prod_{i=1}^n h_i) \tag{23}$$

$$D = B_n \cdot h_0^{\prod_{i=1}^n e_i} \tag{24}$$

and

$$F_j = \prod_{i=1}^n (w_{i,j})^{e_i} \text{ for } j=1..\text{width} \tag{25}$$

Then the following identities are checked, and if all are true, the verification succeeds:

$$A^v \cdot A' = g^{k_A} \prod_{i=1}^n h_i^{k_{E,i}} \quad (26)$$

$$B_i^v \cdot B'_i = g^{k_{B,i}} (B_{i-1})^{k_{E,i}}, \text{ where } B_0 = h_0, \ i = 1..n \quad (27)$$

$$C^v \cdot C' = g^{k_C} \quad (28)$$

$$D^v \cdot D' = g^{k_D} \quad (29)$$

and

$$F_j^v F'_j = Encr(0, -k_{F,j}) \prod_{i=1}^n (w'_{i,j})^{k_{E,i}} \text{ for } j=1..\text{width} \quad (30)$$

3 Performance

Environments used for measuring times:

Workstation

- Ubuntu 22.04.3
- HP Z440 Workstation, Intel Xeon CPU E5-1650 v3 @ 3.50GHz
- 6-cores, two threads per core.

Server

- Ubuntu 22.04.3
- HP Z840 Workstation, Intel Xeon CPU E5-2680 v3 @ 2.50GHz
- 24-cores, two threads per core.

Laptop

- Windows 10 Pro
- Dell Precision M3800, Intel i7-4712HQ CPU @ 2.30GHz
- 4-cores, two threads per core.

3.1 Operation counts

- n = number of rows, eg ballots or contests
- $width$ = number of ciphertexts per row
- $N = nrows * width$ = total number of ciphertexts to be mixed

	shuffle	proof	verify
regular exps	0	$2N + n - 1$	$4N + 4n + 1$
accelerated exps	$2N$	$6n + 2width + 4$	$2n + 2width + 6$

Table 1: Exponent operation count

3.2 Regular vs accelerated exponentiation time

When the same base is used many times for exponentiation, it can be optimized with Pereira’s ”pow-radix” precomputation and Montgomery forms, replacing modular exponentiation with table lookups and multiplies. See [11] for more details. Here we measure the performance difference between regular and accelerated exponentiation, after the acceleration cache warms up:

Workstation

```
acc took 12551 msec for 20000 = 0.62755 msec per acc
exp took 40998 msec for 20000 = 2.0499 msec per exp
exp/acc took 3.266
```

Server

```
acc took 15288 msec for 20000 = 0.7644 msec per acc
exp took 46018 msec for 20000 = 2.3009 msec per exp
exp/acc = 3.010
```

Laptop

```
acc took 16910 msec for 20000 = 0.8455 msec per acc
exp took 55654 msec for 20000 = 2.7827 msec per exp
exp/acc took 3.291
```

From which we can conclude that regular exponentiation is about 3 times slower after the acceleration cache warms up.

We can estimate the performance gain of this acceleration using the operation counts from above, and the factor of 3 for the acceleration speedup. Then

$$speedup = noacc/withacc = totalcount/(expcount + (acccount/3))$$

When width = 34, n rows greater than 50, ShuffleAndProof gets around a 50 percent speedup, while Verify gets only a few percent.

3.3 Timing results

shuffle & proof, Verificatum vs egk-mixnet, HPZ440




Figure 1: compare ShuffleAndProof

verify, Verificatum vs egk-mixnet, HPZ440


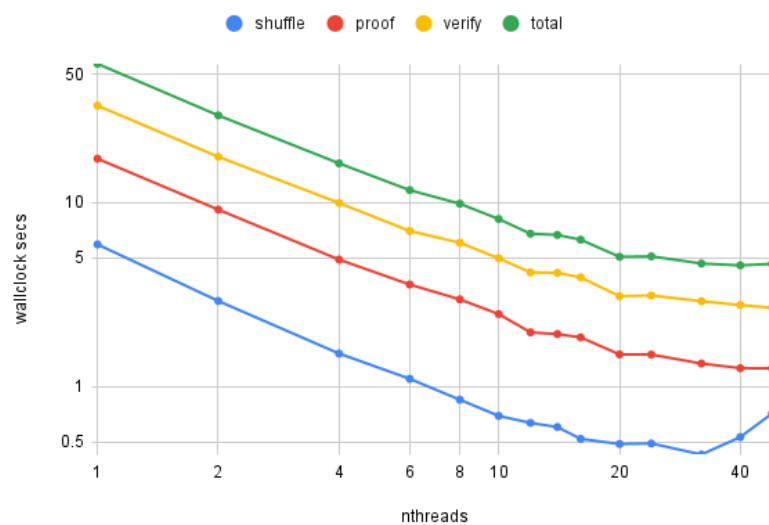


Figure 2: compare Verify

HP Z840 24 cores: 100x34



Dell 4 cores, 100x34

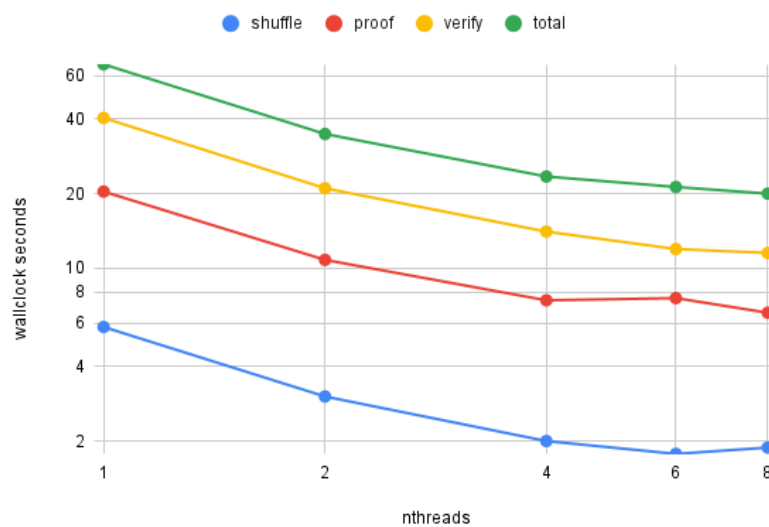


Figure 3: Egk-mixnet time, vary nthreads

3.4 Discussion of Timing results

3.4.1 Parallelization

Egk-mixnet uses Kotlin coroutines to parallelize all operations that can operate independently on row or column vectors or matrices. This parallelization is done at a high level within the algorithm code, so as to put as much work into each thread as possible. Verificatum appears to place its parallelization code lower down in the stack, on individual vector operations.

Figure 1 and 2 above compare wallclock time of egk-mixnet to Verificatum Java (vmn-java), and to Verificatum using the GNU Multiple Precision Arithmetic (GMP) library and the GMP Modular Exponentiation Extension (GMPMEE) library (vmn-gmp). The tests are done on a 6-core HPZ440 workstation, but they are similar on the other two machines. Both tests are for 100 ballots of width 34.

Figure 1 shows times for performing both the Shuffle and the Proof. Egk-mixnet has good parallel speedup up to at least the number of cores, while Verificatum achieves much less parallel speedup. Vmn-gmp is within 20% of Egk when $nthreads = 1$, but then rapidly declines in relative performance as the thread count increases.

Figure 2 for Verify shows much better parallelization on both variants of Verificatum. Vmn-gmp is almost 3x better than egk-mixnet for $nthreads = 1$, but egk-mixnet comes to within 20% of vmn-gmp for $nthreads = 6$, due to better parallelization.

Figure 3 shows the parallelization times for egk-mixnet on the other two machines. Speedup is mostly good up to the number of cores, with some interference from other processes as the high end of that. Shuffle in particular shows the performance falloff when one tries to use all of the virtual threads. The percent of unparallelized code can be estimated from the formula $time = (unparallel) + (parallel)/nthreads$, running with various values of $nthreads$, and solving the linear equations. A rough estimate from the Server machine (100 ballots of width 34) is:

shuffle	proof	verify
1.6 %	4.2 %	4.7 %

Table 2: Percent unparallelized code in egk-miknet

3.4.2 Algorithm differences

The results for `nthreads = 1` measure raw speed without parallel speedup effects. `Vmn-gmp` has a strong advantage over both Java versions. This is due to the use of the `GMPEE` library function `spown()`, which implements a fixed-exponent exponentiation algorithm. This algorithm is used once in the proof (eq 15) and twice in the verify (eq 25 and 30), which are the source of the $2N$ and $4N$ operation counts, respectively. `Vmn-gmp` is faster than `vmn-java` by a factor of approx 2 and 5, respectively. `Vmn-gmp` is slower than `egk-mixnet` by 18% for `ShuffleAndProof`, and faster by 2.7 for `Verify`.

`Egk-mixnet` has a 2.5x speedup over `vmn-java` for `ShuffleAndProof`, and 1.8x speedup for `Verify`. We would expect `egk-mixnet` use of accelerated exponentiation (which `vmn-java` does not have but `vmn-gmp` does) to account for approx 50% speedup for `ShuffleAndProof` and almost nothing for `Verify`. This suggests that besides parallelization speedup, `egk-mixnet` implementation is somehow faster by approx 2x. This is somewhat surprising and bears further investigation.

A ElGamal Encryption and Reencryption

(1.2a)

$$Encr(m, \xi) = (g^\xi, K^{m+\xi}) = (a, b)$$

$$Encr(0, \xi') = (g^{\xi'}, K^{\xi'})$$

(1.2b)

$$(a, b) * (a', b') = (a * a', b * b')$$

$$Encr(m, \xi) * Encr(m', \xi') = (g^{\xi+\xi'}, K^{m+m'+\xi+\xi'}) = Encr(m+m', \xi+\xi')$$

(1.2c)

$$(a, b)^k = (a^k, b^k)$$

$$Encr(m, \xi)^k = (g^{\xi*k}, K^{(m*k+\xi*k)}) = Encr(m*k, \xi*k)$$

(1.2d)

$$\prod_{j=1}^n Encr(m_j, \xi_j) = (g^{\sum_{j=1}^n \xi_j}, K^{\sum_{j=1}^n m_j + \sum_{j=1}^n \xi_j}) = Encr(\sum_{j=1}^n m_j, \sum_{j=1}^n \xi_j)$$

$$\prod_{j=1}^n Encr(m_j, \xi_j)^{k_j} = Encr(\sum_{j=1}^n (m_j * k_j), \sum_{j=1}^n (\xi_j * k_j))$$

(1.2e)

$$ReEncr(m, r) = (g^{\xi+r}, K^{m+\xi+r}) = Encr(0, r) * Encr(m, \xi)$$

$$ReEncr(m, r)^k = Encr(0, r * k) * Encr(m * k, \xi * k)$$

(1.2f)

$$\prod_{j=1}^n ReEncr(e_j, r_j) = (g^{\sum_{j=1}^n (\xi_j + r_j)}, K^{\sum_{j=1}^n (m_j + \xi_j + r_j)})$$

$$= ReEncr(\prod_{j=1}^n e_j, \sum_{j=1}^n r_j)$$

(1.2g)

$$\begin{aligned}
\prod_{j=1}^n ReEncr(m_j, r_j)^{k_j} &= \prod_{j=1}^n Encr(0, r_j * k_j) * \prod_{j=1}^n Encr(m_j * k_j, \xi_j * k_j) \\
&= Encr(0, \sum_{j=1}^n (r_j * k_j)) * \prod_{j=1}^n Encr(m_j, \xi_j)^{k_j}
\end{aligned}$$

Let

- $e_j = Encr(m_j, \xi_j)$
- $re_j = ReEncr(m_j, r_j) = ReEncr(e_j, r_j) = Encr(0, r_j) * e_j$

then

$$\begin{aligned}
re_j &= Encr(0, r_j) * e_j \\
\prod_{j=1}^n re_j^{k_j} &= \prod_{j=1}^n Encr(0, r_j)^{k_j} * \prod_{j=1}^n e_j^{k_j} \\
&= Encr(0, \sum_{j=1}^n (r_j * k_j)) * \prod_{j=1}^n e_j^{k_j}, \text{ (Equation 1)}
\end{aligned}$$

B Alternative Calculation of B and B'

The calculation of B (using e instead of e' here for notational simplicity):

$$\begin{aligned}
B_1 &= g^{b_1} h_0^{e_1} \\
B_i &= g^{b_i} (B_{i-1})^{e_i}, \quad i = 2..n
\end{aligned}$$

can be done to only use accelerated exponents.

Expand the series:

$$\begin{aligned}
B_1 &= g^{b_1} (h_0)^{e_1} \\
B_2 &= g^{b_2} (B_1)^{e_2} = g^{b_2 + b_1 e_2} \cdot h_0^{e_1 \cdot e_2} \\
B_3 &= g^{b_3} (B_2)^{e_3} = g^{b_3 + (b_2 + b_1 e_2) e_3} \cdot h_0^{e_1 \cdot e_2 \cdot e_3} \\
&\dots \\
B_i &= g^{b_i} (B_{i-1})^{e_i} = g^{geps_i} \cdot h_0^{heps_i}
\end{aligned}$$

where

$$\begin{aligned}
geps_1 &= b_1 \\
geps_i &= b_i + (geps_{i-1}) \cdot e_i, \quad i > 1 \\
heps_i &= \prod_{j=1}^i e_j
\end{aligned}$$

Then each row has exactly 2 accelerated exponentiations.

Similarly for B' :

$$\begin{aligned}
B'_1 &= g^{\beta_1} h_0^{eps_1} \\
B'_i &= g^{\beta_i} (B'_{i-1})^{eps_i}, \quad i = 2..n
\end{aligned}$$

can be done with only accelerated exponentiations:

$$B'_i = g^{gpeps_i} \cdot h_0^{hpeps_i}$$

where

$$\begin{aligned}
gpeps_1 &= \beta_1 \\
gpeps_i &= \beta_i + (gpeps_{i-1}) \cdot eps_i, \quad i > 1 \\
hpeps_1 &= eps_1 \\
hpeps_i &= hpeps_{i-1} \cdot eps_i, \quad i > 1
\end{aligned}$$

The net result is that this shifts 2n operations from exp to acc in the proof.

C Implementation notes

Permutation The permutation function is defined in Verificatum as the inverse of the definition here, and the operations of permute and invert are therefore switched. This is really a notational difference, and does not affect the mathematics. Its also worth noting that only the Shuffle/Prover works with the permutation, and the Verifier is not affected by the switch.

Independent Generators Verificatum [3] gives warnings in section 6.8 and section 8.2 on choosing the generators. It also uses a complex algorithm involving “statistical error”. These need to be evaluated.

References

- [1] B. Terelius and D. Wikström. *Proofs of restricted shuffles*, In D. J. Bernstein and T. Lange, editors, AFRICACRYPT’10, 3rd International Conference on Cryptology in Africa, LNCS 6055, pages 100–113, Stellenbosch, South Africa, 2010.
- [2] D. Wikström. *A commitment-consistent proof of a shuffle*. In C. Boyd and J. González Nieto, editors, ACISP’09, 14th Australasian Conference on Information Security and Privacy, LNCS 5594, pages 407–421, Brisbane, Australia, 2009.
- [3] D. Wikström. *How to Implement a Stand-alone Verifier for the Verificatum Mix-Net VMN Version 3.1.0*, 2022-09-10, <https://www.verificatum.org/files/vmnv-3.1.0.pdf>.
- [4] D. Wikström. *Verificatum Mix-Net*, <https://github.com/verificatum/verificatum-vmn>.
- [5] Rolf Haenni, Reto E. Koenig, Philipp Locher, Eric Dubuis, *CHVote Protocol Specification Version 3.5*, Bern University of Applied Sciences, February 28th, 2023, <https://eprint.iacr.org/2017/325.pdf>
- [6] R. Haenni, P. Locher, R. E. Koenig, and E. Dubuis, *Pseudo-code algorithms for verifiable re-encryption mix-nets*, In M. Brenner, K. Rohloff, J. Bonneau, A. Miller, P. Y. A. Ryan, V. Teague, A. Bracciali, M. Sala, F. Pintore, and M. Jakobsson, editors, FC’17, 21st International

Conference on Financial Cryptography, LNCS 10323, pages 370–384, Silema, Malta, 2017.

- [7] *OpenCHVote* E-Voting Group, Institute for Cybersecurity and Engineering, Bern University of Applied Sciences, <https://gitlab.com/openchvote/cryptographic-protocol>
- [8] Thomas Haines, *A Description and Proof of a Generalised and Optimised Variant of Wikström’s Mixnet*, arXiv:1901.08371v1 [cs.CR], 24 Jan 2019.
- [9] Josh Benaloh and Michael Naehrig, *ElectionGuard Design Specification, Version 2.0.0*, Microsoft Research, August 18, 2023, https://github.com/microsoft/electionguard/releases/download/v2.0/EG_Spec_2_0.pdf.
- [10] John Caron, Dan Wallach, *ElectionGuard Kotlin library*, <https://github.com/votingworks/electionguard-kotlin-multiplatform>.
- [11] Dan Wallach, *ElectionGuard Kotlin cryptography notes*, <https://github.com/votingworks/electionguard-kotlin-multiplatform/blob/main/docs/CryptographyNotes.md>
- [12] John Caron, *egk-mixnet library*, <https://github.com/JohnLCaron/egk-mixnet>.